


# Machine Learning

## 4771


Instructor: Itsik Pe'er

# Reminder: Gaussian Noise




# Regression

## Classification





## Regression, $f(x)=y$




Supervised

## Density/Structure Estimation




## Clustering




Unsupervised

## Feature Selection



## Anomaly Detection



# Outline

- Regression
- Empirical Risk Minimization
- Least Squares
- Higher Order Polynomials
- Under-fitting / Over-fitting
- Cross-Validation

# Function Approximation

- Start with training dataset

$$X = \{(x_1, y_1), (x_1, y_1), \dots, (x_N, y_N)\} \quad x \in \mathbf{R}^D = \begin{bmatrix} x(1) \\ x(2) \\ \vdots \\ x(D) \end{bmatrix}, y \in \mathbf{R}$$


- Have  $N$  (input, output ) pairs
- Find a function  $f(x)$  to predict  $y$  from  $x$   
That fits the training data well

# Function Approximation

- Start with training dataset

$$X = \{(x_1, y_1), (x_1, y_1), \dots, (x_N, y_N)\} \quad x \in \mathbf{R}^D = \begin{bmatrix} x(1) \\ x(2) \\ \vdots \\ x(D) \end{bmatrix}, y \in \mathbf{R}$$

- Have  $N$  (input, output) pairs
- Find a function  $f(x)$  to predict  $y$  from  $x$   
That fits the training data well




- Example: predict your grade  $y$  in ML  
 $x = [\text{HW grades}; \text{quiz grades}; \text{midterm}; \text{final}]$

# Function Approximation

- Start with training dataset

$$X = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\} \quad x \in \mathbf{R}^D = \begin{bmatrix} x(1) \\ x(2) \\ \vdots \\ x(D) \end{bmatrix}, y \in \mathbf{R}$$

- Have  $N$  (input, output) pairs
- Find a function  $f(x)$  to predict  $y$  from  $x$   
That fits the training data well




- Example: predict your grade  $y$  in ML  
 $x = [\text{HW grades; quiz grades; midterm; final}]$   
 $x = [\text{HW 1-4 grades; quiz 1 grade; midterm}]$

# Function Approximation

- Start with training dataset

$$X = \{(x_1, y_1), (x_1, y_1), \dots, (x_N, y_N)\} \quad x \in \mathbf{R}^D = \begin{bmatrix} x(1) \\ x(2) \\ \vdots \\ x(D) \end{bmatrix}, y \in \mathbf{R}$$

- Have N (input, output ) pairs
- Find a function  $f(x)$  to predict  $y$  from  $x$   
That fits the training data well



- Example: predict the price of house in dollars  $y$  using  $x = [\text{\#rooms}; \text{latitude}; \text{longitude}; \dots]$
- Need:
  - a) Way to evaluate how good a fit we have
  - b) Class of functions in which to search for  $f(x)$



# Empirical Risk Minimization

- Idea: minimize 'loss' on the training data set
- Empirical = use the training set to find the best fit
- Define a loss function of how good we fit a single point:

$$Loss(y, f(x)) = -\log Prob(y|f(x))$$


# Empirical Risk Minimization

- Idea: minimize 'loss' on the training data set
- Empirical = use the training set to find the best fit
- Define a loss function of how good we fit a single point:

$$Loss(y, f(x)) = -\log Prob(y|f(x))$$

- Empirical Risk = the average loss over the dataset

$$R = \frac{1}{N} \sum_{i=1}^N Loss(y_i, f(x_i))$$



# Empirical Risk Minimization

- Idea: minimize 'loss' on the training data set
- Empirical = use the training set to find the best fit
- Define a loss function of how good we fit a single point:

$$Loss(y, f(x)) = -\log Prob(y|f(x))$$

- Empirical Risk = the average loss over the dataset



$$R = \frac{1}{N} \sum_{i=1}^N Loss(y_i, f(x_i))$$

- Gaussian noise: squared error from y value

$$Loss(y_i, f(x_i)) = \frac{1}{2} (y_i - f(x_i))^2$$

- Other possible loss: absolute error

$$Loss(y_i, f(x_i)) = |y_i - f(x_i)|$$




# Linear Function Classes

- Linear is simplest class of functions to search over:

$$f(x; \theta) = \theta^T x + \theta_0 = \theta_0 + \sum_{d=1}^D \theta_d x(d)$$

- Start with  $x$  being 1-dimensional ( $D=1$ ):

$$f(x; \theta) = \theta_1 x + \theta_0$$



# Linear Function Classes



- Linear is simplest class of functions to search over:

$$f(x; \theta) = \theta^T x + \theta_0 = \theta_0 + \sum_{d=1}^D \theta_d x(d)$$

- Start with  $x$  being 1-dimensional ( $D=1$ ):

$$f(x; \theta) = \theta_1 x + \theta_0$$

- Plug in the above & minimize empirical risk over  $\theta$



$$R(\theta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)^2$$

- Note: minimum occurs when  $R(\theta)$  gets flat (not always!)
- Note: when  $R(\theta)$  is flat, gradient  $\nabla_{\theta} R = 0$

# Min by Gradient=0

- Gradient=0 means the partial derivatives are all 0
- Take partials of empirical risk:

$$R(\theta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)^2$$

$$\nabla_{\theta} R = \begin{bmatrix} \frac{\partial R}{\partial \theta_0} \\ \frac{\partial R}{\partial \theta_1} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

# Min by Gradient=0

- Gradient=0 means the partial derivatives are all 0
- Take partials of empirical risk:

$$R(\theta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)^2$$

$$\frac{\partial R}{\partial \theta_0} = \frac{1}{N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)(-1) = 0$$


$$\frac{\partial R}{\partial \theta_1} = \frac{1}{N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)(-x_i) = 0$$

$$\nabla_{\theta} R = \begin{bmatrix} \frac{\partial R}{\partial \theta_0} \\ \frac{\partial R}{\partial \theta_1} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

# Interpreting the equations

$$\frac{1}{N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0) = 0$$

$$\frac{1}{N} \sum_{i=1}^N x_i (y_i - \theta_1 x_i - \theta_0) = 0$$






# Interpreting the equations

$$\frac{1}{N} \sum_{i=1}^N \overbrace{(y_i - \theta_1 x_i - \theta_0)}^{err_i} = 0$$

$$\hat{E}(Err) = Avg(err_i) = \frac{1}{N} \sum_{i=1}^N err_i = 0$$

$$\frac{1}{N} \sum_{i=1}^N x_i (y_i - \theta_1 x_i - \theta_0) = 0$$

$$\widehat{Cov}(X, Err) = \hat{E}(x_i err_i) = \frac{1}{N} \sum_{i=1}^N x_i err_i = 0$$



# Solving Least Squares

$$\frac{1}{N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0) = 0$$

$$\frac{1}{N} \sum_{i=1}^N x_i (y_i - \theta_1 x_i - \theta_0) = 0$$

# Solving Least Squares

$$\frac{1}{N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0) = 0$$

$$\frac{1}{N} \sum_{i=1}^N x_i (y_i - \theta_1 x_i - \theta_0) = 0$$

$$\theta_0 = \frac{1}{N} \sum_{i=1}^N y_i - \theta_1 \frac{1}{N} \sum_{i=1}^N x_i$$

$$\theta_1 \sum_{i=1}^N x_i^2 = \sum_{i=1}^N x_i y_i - \theta_0 \sum_{i=1}^N x_i$$

$$\theta_1 = \frac{\sum x_i y_i - \frac{1}{N} \sum y_i \sum x_i}{\sum x_i^2 - \frac{1}{N} \sum x_i \sum x_i}$$

# Multi-Dimensional Regression

- More elegant/general to solve  $\nabla_{\theta} R = 0$  with linear algebra
- Rewrite empirical risk in vector-matrix notation:

$$R(\theta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)^2$$

=

Can add more dimensions by adding columns to  $X$  matrix and rows to  $\theta$  vector

# Multi-Dimensional Regression

- More elegant/general to solve  $\nabla_{\theta} R = 0$  with linear algebra
- Rewrite empirical risk in vector-matrix notation:

$$\begin{aligned} R(\theta) &= \frac{1}{2N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)^2 \\ &= \frac{1}{2N} \sum_{i=1}^N \left( y_i - [1 \quad x_i] \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} \right)^2 \\ &= \frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} - \begin{bmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} \right\|^2 \\ &= \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\theta\|^2 \end{aligned}$$

# Multi-Dimensional Regression

- More elegant/general to solve  $\nabla_{\theta} R = 0$  with linear algebra
- Rewrite empirical risk in vector-matrix notation:

$$\begin{aligned} R(\theta) &= \frac{1}{2N} \sum_{i=1}^N (y_i - \theta_1 x_i - \theta_0)^2 \\ &= \frac{1}{2N} \sum_{i=1}^N \left( y_i - [1 \quad x_i] \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} \right)^2 \\ &= \frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} - \begin{bmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} \right\|^2 \\ &= \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\theta\|^2 \end{aligned}$$

Can add more dimensions by adding columns to  $\mathbf{X}$  matrix and rows to  $\theta$  vector

# Multi-Dimensional Regression

- More elegant/general to solve  $\nabla_{\theta} R = 0$  with linear algebra
- Rewrite empirical risk in vector-matrix notation:

$$R(\theta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \theta_1 x_i(1) - \dots - \theta_D x_i(D) - \theta_0)^2$$

$$= \frac{1}{2N} \sum_{i=1}^N \left( y_i - [1 \quad x_i(1) \dots x_i(D)] \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_D \end{bmatrix} \right)^2$$

$$= \frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} - \begin{bmatrix} 1 & x_1(1) \dots x_1(D) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_N(1) \dots x_N(D) \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_D \end{bmatrix} \right\|^2$$

$$= \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\theta\|^2$$

Can add more dimensions by adding columns to  $\mathbf{X}$  matrix and rows to  $\theta$  vector

# Multi-Dimensional Regression

- More realistic dataset: many measurements
- Have  $N$  apartments each with  $D$  measurements
- Each row of  $X$  is [#rooms; latitude; longitude,...]

$$X = \begin{bmatrix} 1 & x_1(1) \cdots & x_1(D) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_N(1) \cdots & x_N(D) \end{bmatrix}$$



	<b>1212 Fifth Avenue PENTHOUSE</b> Condo, Upper Carnegie Hill Listed by <b>Nancy Packes Inc.</b>	<b>\$7,995,000</b> 3 beds 3.5 baths 2,689 ft <sup>2</sup>
	<b>210 East 73rd Street #PHB</b> Co-op, Upper East Side Listed by <b>Brown Harris Stevens</b>	<b>\$3,495,000</b> 2 beds 3 baths
	<b>66 East 11th Street</b> Building, Greenwich Village Listed by <b>Douglas Elliman</b>	<b>\$120,000,000</b>
	<b>150 West 56th Street #PH</b> Condo, Midtown Listed by <b>Douglas Elliman</b>	<b>\$100,000,000</b> 6 beds 9 baths 8,000 ft <sup>2</sup>
	<b>50 Central Park South #PH34/35</b> Condo, Central Park South Listed by <b>Halstead Property</b>	<b>\$95,000,000</b> 3 beds 3.5 baths
	<b>15 Central Park West #35S</b> Condo, Lincoln Square Listed by <b>CORE</b>	<b>\$95,000,000</b> 5 beds 5+ baths
	<b>828 Fifth Avenue #XXX</b> Co-op, Lenox Hill Listed by <b>Stribling</b>	<b>\$72,000,000</b> 8 beds 10.5 baths
	<b>785 Fifth Avenue #PH1718</b> Co-op, Lenox Hill Listed by <b>Corcoran</b>	<b>\$65,000,000</b> <b>IN CONTRACT</b> 7 beds 11 baths



# Multi-Dimensional Regression

- Solving gradient=0  $\nabla_{\theta} R = 0$

$$\nabla_{\theta} \left( \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\theta\|^2 \right) = 0$$

# Multi-Dimensional Regression

- Solving gradient=0  $\nabla_{\theta} R = 0$

$$\nabla_{\theta} \left( \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\theta\|^2 \right) = 0$$

$$\frac{1}{2N} \nabla_{\theta} ((\mathbf{y} - \mathbf{X}\theta)^T (\mathbf{y} - \mathbf{X}\theta)) = 0$$

$$\frac{1}{2N} \nabla_{\theta} (\mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}\theta + \theta^T \mathbf{X}^T \mathbf{X}\theta) = 0$$

$$\frac{1}{2N} (-2\mathbf{X}^T \mathbf{y} + 2\mathbf{X}^T \mathbf{X}\theta) = 0$$

$$\mathbf{X}^T \mathbf{X}\theta = \mathbf{X}^T \mathbf{y}$$

$$\theta^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

# Multi-Dimensional Regression

- Solving gradient=0

$$\begin{aligned} X^T X \theta &= X^T y \\ \theta^* &= (X^T X)^{-1} X^T y \end{aligned}$$


- If the matrix  $X$  is skinny, the solution is probably unique
- If  $X$  is fat (more dimensions than points) we get multiple solutions for theta which give zero error.
- The pseudoinverse (`numpy.linalg.pinv(X)`) returns the theta with zero error and which has the smallest norm.

$$\min_{\theta} \|\theta\|^2 \text{ such that } X\theta = y$$

# 2D Linear Regression

- Once best  $\theta^*$  is found, we can plug it into the function:

$$f(x; \theta) = \theta_2^* x(2) + \theta_1^* x(1) + \theta_0^*$$




- What would a fat  $X$  look like?

# Polynomial Function Classes

- Back to 1-dim  $x$  ( $D=1$ ) BUT Nonlinear

- Polynomial: 
$$f(x; \theta) = \sum_{p=0}^P \theta_p x^p$$




# Polynomial Function Classes

- Back to 1-dim  $x$  ( $D=1$ ) BUT Nonlinear

- Polynomial: 
$$f(x; \theta) = \sum_{p=0}^P \theta_p x^p$$

- Writing Risk:

$$R(\theta) = \frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} - \begin{bmatrix} 1 & x_1^1 \cdots & x_1^P \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_N^1 \cdots & x_N^P \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_P \end{bmatrix} \right\|^2$$






- Order- $P$  polynomial regression fitting for 1D variable is same as  $P$ -dimensional linear regression!

- Construct a multidim  $\mathbf{X}_i = [x_i^0 \quad x_i^1 \quad x_i^2 \quad x_i^3]$   
 $\mathbf{X}$ -vector from  $x$  scalar

- More generally any  $\mathbf{X}_i = [\phi_0(x_i) \quad \phi_1(x_i) \quad \phi_2(x_i) \quad \phi_3(x_i)]$







# Underfitting/Overfitting

- Try varying  $P$ . Higher  $P$  fits a more complex function class
- Observe  $R(\theta^*)$  drops with bigger  $P$



# Underfitting/Overfitting


- Try varying  $P$ . Higher  $P$  fits a more complex function class
- Observe  $R(\theta^*)$  drops with bigger  $P$






# Evaluating The Regression

- Unfair to use empirical to find best order  $P$
- High  $P$  (vs.  $N$ ) can overfit, even linear case!
- $\min R(\theta^*)$  not on training but on future data
- Want model to *Generalize* to future data



# Evaluating The Regression

- Unfair to use empirical to find best order  $P$
- High  $P$  (vs.  $N$ ) can overfit, even linear case!
- $\min R(\theta^*)$  not on training but on future data
- Want model to *Generalize* to future data



True loss:  $R_{true}(\theta) = \int p(x, y) \frac{1}{2} (y - \theta^T x)^2 dx dy$

- One approach: split data into training / testing portion


$$\{(x_1, y_1), \dots, (x_N, y_N)\} \quad \{(x_{N+1}, y_{N+1}), \dots, (x_{N+M}, y_{N+M})\}$$

- Estimate  $\theta^*$  with training loss:  $R_{train}(\theta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \theta^T x_i)^2$

- Evaluate  $P$  with testing loss:  $R_{test}(\theta^*) = \frac{1}{2M} \sum_{i=N+1}^{N+M} (y_i - \theta^{*T} x_i)^2$

# Crossvalidation

- Try fitting with different polynomial order  $P$
- Select  $P$  which gives lowest  $R_{test}(\theta^*)$



- Think of  $P$  as a measure of the complexity of the model
- Higher order polynomials are more flexible and complex

# Summary

- ◆ Regression: approximating a function
- ◆ Minimizing loss function
- ◆ Least squares solution to linear regression
- ◆ Vector/matrix representation
- ◆ Hi-dimension, polynomial regression
- ◆ Overfit/underfit