

Machine Learning

4771

Instructor: Itsik Pe'er

Administration

◆ Quiz:

- 30 minutes
- Multiple choice
- Mudd 833/PUP428/CVN proctorship/IDS

◆ 2nd Quiz: April 10th


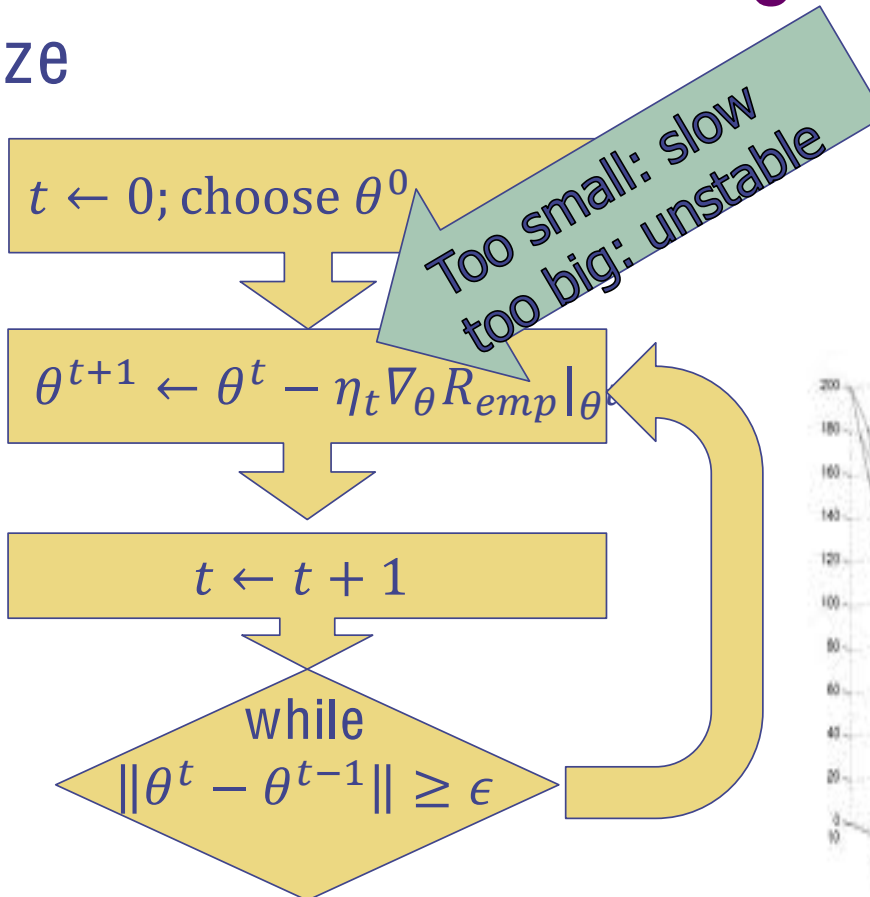
Piazza

- ◆ A professional, not social forum
- ◆ Avoid offensive language
- ◆ Report issues to me/head TA
- ◆ Avoid staff feedback
 - Send to me, iachair, CULPA, class evaluation
- ◆ No bullying

- ◆ Violators will be kicked off the forum

Gradient Descent Algorithm


Initialize



Class 7

- Perceptrons
- Online & Stochastic Gradient Descent
- Convergence Guarantee
- Gap tolerance

Perceptron (another Neuron)



Perceptron (another Neuron)


- Classification scenario once again but consider +1, -1 labels

$$X = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}, x \in \mathbf{R}^D, y \in \{-1, 1\}$$




- A better choice for a classification squashing function is

Loss class g, f

$$g(z) = \begin{cases} -1 & \text{when } z < 0 \\ +1 & \text{when } z \geq 0 \end{cases}$$




$y = f$ $y \neq f$
 $yf = 1$ $yf = -1$




Perceptron (another Neuron)

- A better choice for a classification squashing function is

$$g(z) = \begin{cases} -1 & \text{when } z < 0 \\ +1 & \text{when } z \geq 0 \end{cases}$$




- And a better choice is classification loss



Perceptron (another Neuron)

- A better choice for a classification squashing function is

$$g(z) = \begin{cases} -1 & \text{when } z < 0 \\ +1 & \text{when } z \geq 0 \end{cases}$$




- And a better choice is classification loss

$$\begin{aligned} \text{Loss}^{\text{class}}(y, f(\mathbf{x}; \theta)) &= \text{step}(-yf(\mathbf{x}; \theta)) \\ \text{step}(z) &= \begin{cases} 1 & z > 0 \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

- What does this $R(\theta)$ function look like?

$$R^{\text{class}}(\theta) = \frac{1}{N} \sum_{i=1}^N \text{step}(-y_i \theta^T \mathbf{x}_i) = \frac{1}{4N} \sum_{i=1}^N (y_i - g(\theta^T \mathbf{x}_i))^2$$




Perceptron & Classification Loss

- Classification loss for the Risk leads to hard minimization
- What does this $R(\theta)$ function look like?

$$R^{class}(\theta) = \frac{1}{N} \sum_{i=1}^N \text{step}(-y_i \theta^T x_i)$$

- Can't do gradient descent since the gradient is zero except at edges when a label flips





Perceptron & Perceptron Loss

- Instead of Classification Loss

$$R^{class}(\theta) = \frac{1}{N} \sum_{i=1}^N \text{step}(-y_i \theta^T x_i)$$

- Consider Perceptron Loss:

$$R^{per}(\theta) = \frac{1}{N} \sum_{i \in \text{misclassified}} y_i (\theta^T x_i)$$




$$\nabla_{\theta} R^{per}(\theta) = \frac{1}{N} \sum_{i \in \text{misclassified}} y_i x_i$$

$$\theta^{t+1} \leftarrow \theta^t - \eta \sum_{i \in \text{misclassified}} y_i x_i$$

Perceptron & Perceptron Loss


- Instead of Classification Loss

$$R^{class}(\theta) = \frac{1}{N} \sum_{i=1}^N \text{step}(-y_i \theta^T \mathbf{x}_i)$$



- Consider Perceptron Loss:

$$R^{per}(\theta) = -\frac{1}{N} \sum_{i \in \text{misclassified}} y_i (\theta^T \mathbf{x}_i)$$



- Instead of staircase-shaped R get smooth piece-wise linear
- Get reasonable gradients for gradient descent

$$\nabla_{\theta} R^{per}(\theta) = -\frac{1}{N} \sum_{i \in \text{misclassified}} y_i \mathbf{x}_i$$

$$\theta^{t+1} = \theta^t - \eta \nabla_{\theta} R^{per} \Big|_{\theta^t} = \theta^t + \eta \frac{1}{N} \sum_{i \in \text{misclassified}} y_i \mathbf{x}_i$$



Perceptron vs. Linear Regression

- Linear regression gets close but doesn't do perfectly

classification error = 2
squared error = 0.139

- Perceptron gets zero error

classification error = 0
perceptron err = 0



Stochastic Gradient Descent

- Gradient Descent vs. Stochastic Gradient Descent
- Computing average gradient for all points for taking a step

$$\nabla_{\theta} R^{per}(\theta) = -\frac{1}{N} \sum_{i \in \text{misclassified}} y_i \mathbf{x}_i$$

Stochastic Gradient Descent

- Gradient Descent vs. Stochastic Gradient Descent
- Instead of computing the average gradient for all points and then taking a step

$$\nabla_{\theta} R^{per}(\theta) = -\frac{1}{N} \sum_{i \in \text{misclassified}} y_i \mathbf{x}_i$$

- Update the gradient for each mis-classified point by itself


$$\nabla_{\theta} Loss^{per}(\theta) = -y_i \mathbf{x}_i \text{ if } i \text{ misclassified}$$

- Also, set η to 1

$$\theta^{t+1} = \theta^t - \eta \nabla_{\theta} Loss^{per}|_{\theta^t} = \theta^t + y_i \mathbf{x}_i \text{ if } i \text{ misclassified}$$

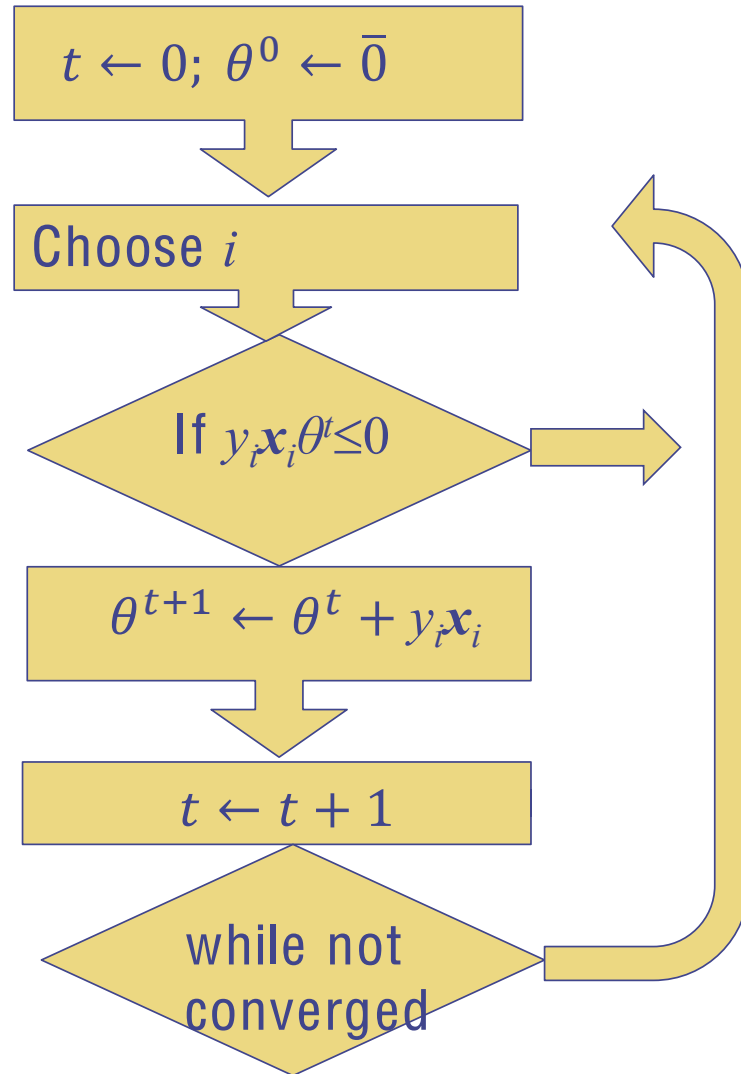
Online Perceptron

- Apply stochastic gradient descent to a perceptron
- Get the “online perceptron” algorithm:



Online Perceptron

- Initialize & repeat



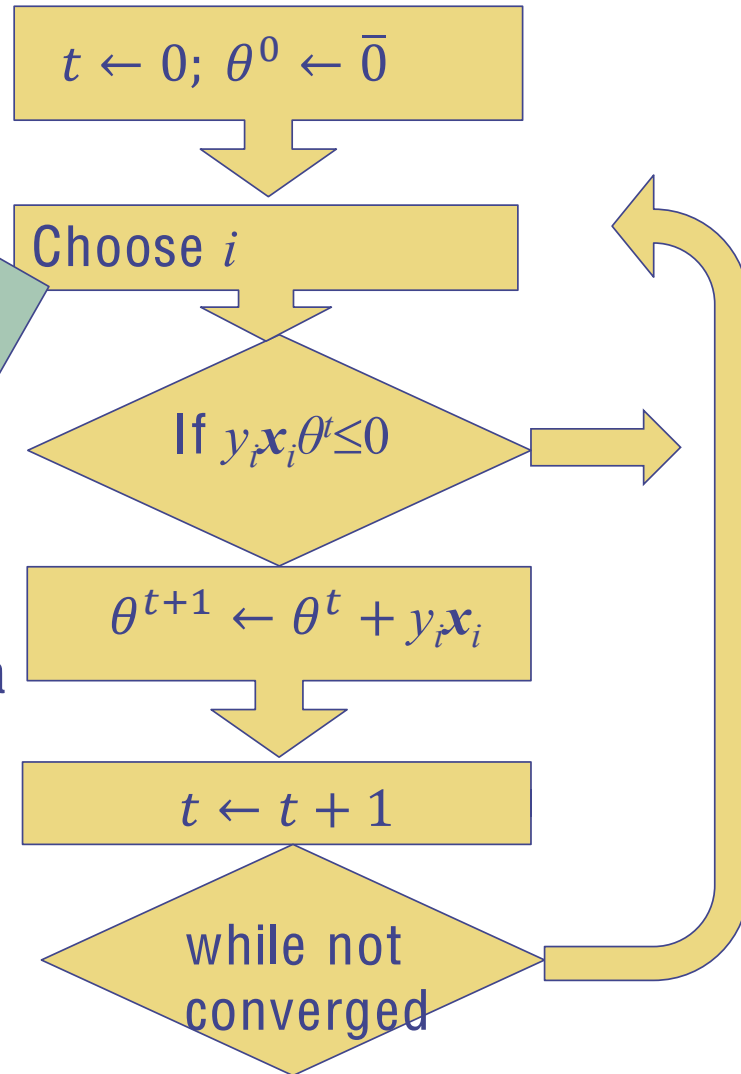
Online Perceptron

- Initialize & repeat

Randomly/iteratively

- If the algorithm stops,
we have θ that separates data

- t = total number of mistakes



Online Perceptron Theorem

Theorem: the online perceptron algorithm converges to zero error in finite t if we assume

1) all data inside a sphere of radius r : $\|x_i\| \leq r \forall i$

2) data is separable with margin γ : $y_i(\theta^*)^T x_i \geq \gamma \forall i$

$$\begin{aligned}
 \underline{\theta^{*T} \theta^t} &= \theta^{*T} (\theta^{t-1} + y_i x_i) = \\
 &= \theta^{*T} \theta^{t-1} + \underline{y \theta^{*T} x_i} \geq \underline{\theta^{*T} \theta^{t-1}} + \underline{\gamma} \\
 &\geq \gamma t
 \end{aligned}$$

Online Perceptron Theorem

Theorem: the online perceptron algorithm converges to zero error in finite t if we assume

- 1) all data inside a sphere of radius r : $\|x_i\| \leq r \forall i$
- 2) data is separable with margin γ : $y_i(\theta^*)^T x_i \geq \gamma \forall i$

Proof:

- Part 1) Look at inner product of current θ^t with θ^*
assume we just updated a mistake on point i :

$$(\theta^*)^T \theta^t = (\theta^*)^T \theta^{t-1} + y_i(\theta^*)^T x_i \geq (\theta^*)^T \theta^{t-1} + \gamma$$

after applying t such updates, we must get:

$$(\theta^*)^T \theta^t \geq t\gamma$$

Online Perceptron Proof

•Part 1) $(\theta^*)^T \theta^t \geq t\gamma$

•Part 2) $\underline{\|\theta^t\|^2} = \left\| \theta^{t-1} + \cancel{y_i x_i} \right\|^2$ ↖ r^2

$$= \underline{\|\theta^{t-1}\|^2} + 2y_i \theta^{t-1 \top} x_i + \|x_i\|^2$$

$$\leq \|\theta^{t-1}\|^2 + r^2 \stackrel{\sim 0}{\leq} t r^2$$

Online Perceptron Proof

- Part 1) $(\theta^*)^T \theta^t \geq t\gamma$
- Part 2) $\|\theta^t\|^2 = \|\theta^{t-1} + y_i \mathbf{x}_i\|^2 = \|\theta^{t-1}\|^2 + 2y_i(\theta^{t-1})^T \mathbf{x}_i + \|\mathbf{x}_i\|^2$
 $\leq \|\theta^{t-1}\|^2 + \|\mathbf{x}_i\|^2 \leq \|\theta^{t-1}\|^2 + r^2 \leq tr^2$

since only update mistakes
middle term is negative

$$|\geq \cos(\theta^*, \theta^t) = \frac{\theta^{*T} \theta^t}{\|\theta^*\| \|\theta^t\|} \geq \frac{t\gamma}{\|\theta^*\| \sqrt{tr^2}}$$

Online Perceptron Proof

- Part 1) $(\theta^*)^T \theta^t \geq t\gamma$
- Part 2) $\|\theta^t\|^2 = \|\theta^{t-1} + y_i \mathbf{x}_i\|^2 = \|\theta^{t-1}\|^2 + 2y_i(\theta^{t-1})^T \mathbf{x}_i + \|\mathbf{x}_i\|^2$
 $\leq \|\theta^{t-1}\|^2 + \|\mathbf{x}_i\|^2 \leq \|\theta^{t-1}\|^2 + r^2 \leq tr^2$

$\cos \leq 1$ since only update mistakes
middle term is negative
- Part 3) Angle between optimal & current solution
$$\cos(\theta^*, \theta^t) = \frac{(\theta^*)^T \theta^t}{\|\theta^t\| \|\theta^*\|} \geq \frac{t\gamma}{\|\theta^t\| \|\theta^*\|} \geq \frac{t\gamma}{\sqrt{tr^2} \|\theta^*\|}$$


apply part 1
then part 2
- Since $\cos \leq 1$, $\frac{t\gamma}{\sqrt{tr^2}} \leq 1$, thus $t \leq \frac{r^2}{\gamma^2} \|\theta^*\|^2$

...so t is finite!

Minimum Training Error?


- Is minimizing Empirical Risk the right thing?
- Are Perceptrons & Neural Networks giving the best classifier?

- Perceptrons are giving a bunch of solutions:



Minimum Training Error?

- Is minimizing Empirical Risk the right thing?
- Are Perceptrons & Neural Networks giving the best classifier?
- We are getting: minimum training error
 not minimum testing error
- Perceptrons are giving a bunch of solutions:



... a better solution → gap tolerant classifier

Empirical Risk Minimization

- Recall linear classifier $f(\mathbf{x}; \theta) = \text{sign}(\theta^T \mathbf{x} + \theta_0) \in \{-1, 1\}$
- Recall ERM: $R_{emp}(\theta) = \frac{1}{N} \sum_i^N \text{Loss}(y_i, f(\mathbf{x}_i; \theta)) \in [0, 1]$
- Some loss functions: quadratic: $\text{Loss}(y, \mathbf{x}, \theta) = \frac{1}{2}(y - f(\mathbf{x}; \theta))^2$
 linear: $\text{Loss}(y, \mathbf{x}, \theta) = |y - f(\mathbf{x}; \theta)|$
 binary: $\text{Loss}(y, \mathbf{x}, \theta) = \text{step}(-yf(\mathbf{x}; \theta))$
- Empirical $R_{emp}(\theta)$ *approximates* the true risk (expected error)

$$R(\theta) = E_P\{\text{Loss}(y, \mathbf{x}, \theta)\} = \int_{\mathbf{X} \times \mathbf{Y}} P(\mathbf{x}, y) \text{Loss}(y, \mathbf{x}, \theta) d\mathbf{x} dy \in [0, 1]$$

Empirical Risk Minimization

- Recall ERM: $R_{emp}(\theta) = \frac{1}{N} \sum_i^N \text{Loss}(y_i, f(x_i; \theta)) \in [0, 1]$
- Empirical $R_{emp}(\theta)$ *approximates* the true risk (expected error)

$$R(\theta) = E_P\{\text{Loss}(y, \mathbf{x}, \theta)\} = \int_{\mathbf{X} \times \mathbf{Y}} P(\mathbf{x}, y) \text{Loss}(y, \mathbf{x}, \theta) d\mathbf{x} dy \in [0, 1]$$
- But, we don't know the true $P(\mathbf{x}, y)$!
- Good news: for any θ , if infinite data, by *law of large numbers*:

$$\lim_{n \rightarrow \infty} R_{emp}(\theta) = R(\theta)$$

- Bad news: ERM may not converge to optimum even if $N \rightarrow \infty$:

$$\text{argmin}_{\theta} R_{emp}(\theta) \neq \text{argmin}_{\theta} R(\theta)$$

...ERM is not consistent

Summary

- Perceptrons:
 - Shoot for perfect classification
 - Optimized by online (stochastic) Gradient Descent
 - Convergence guaranteed
- Gaps required to guard against overfit