

Machine Learning 4771



Instructor: Itsik Pe'er


Reminder: Gaussians Classifiers

Dependent features

ML/Bayesian estimation
of parameters

Mahalanobis distance






Gaussians: Key Points

- If $z \in \mathbf{R}^D$ multivariate normal $z \sim N(\mu, \Sigma)$ (also $z \sim MVN(\mu, \Sigma)$)

$$p(z|\mu, \Sigma) = \frac{1}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}} \exp\left(-\frac{1}{2} (z - \mu)^T \Sigma^{-1} (z - \mu)\right)$$



Gaussians: Key Points

- If $z \in \mathbf{R}^D$ multivariate normal $z \sim N(\mu, \Sigma)$ (also $z \sim MVN(\mu, \Sigma)$)

$$p(z|\mu, \Sigma) = \frac{1}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}} \exp\left(-\frac{1}{2}(z - \mu)^T \Sigma^{-1}(z - \mu)\right)$$

- Affine transformation $w = Az + b \sim MVN(A\mu, A\Sigma A^T)$:

$$\begin{aligned} p_w(w) &\propto p_z(A^{-1}(w - b)) \\ &\propto \exp\left(-\frac{(A^{-1}(w - b) - \mu)^T \Sigma^{-1}(A^{-1}(w - b) - \mu)}{2}\right) \\ &= \exp\left(-\frac{(w - (A\mu + b))^T A^{-1^T} \Sigma^{-1} A^{-1}(w - (A\mu + b))}{2}\right) \end{aligned}$$


Concatenating Gaussians

- Have input and output, each Gaussian:



$$\{(x_1, y_1), \dots, (x_N, y_N)\} \quad x \in \mathbf{R}^{D_x}, y \in \mathbf{R}^{D_y}, D = D_x + D_y$$

$$\text{concatenate } z_i = \begin{bmatrix} x_i \\ y_i \end{bmatrix}$$


$$p(z|\mu, \Sigma) = \frac{1}{(2\pi)^{D/2} \sqrt{|\Sigma|}} \exp \left(-\frac{1}{2} (z - \mu)^T \Sigma^{-1} (z - \mu) \right)$$



Independent:




y




x

Dependent:




Marginals/Conditionals for Independent Gaussians



$$p(x, y) = \frac{\exp\left(-\frac{1}{2}\left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)^T \begin{bmatrix} \Sigma_{xx} & 0 \\ 0 & \Sigma_{yy} \end{bmatrix}^{-1} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)\right)}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}} =$$

Marginals/Conditionals for Independent Gaussians



$$p(x, y) = \frac{\exp\left(-\frac{1}{2}\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)^T \begin{bmatrix} \Sigma_{xx} & 0 \\ 0 & \Sigma_{yy} \end{bmatrix}^{-1} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}} = F_{xx} F_{yy}$$


$$\begin{bmatrix} \Sigma_{xx} & 0 \\ 0 & \Sigma_{yy} \end{bmatrix}^{-1} = \begin{bmatrix} \Sigma_{xx}^{-1} & 0 \\ 0 & \Sigma_{yy}^{-1} \end{bmatrix}$$

$$F_{xx} = \frac{\exp\left(-\frac{(x - \mu_x)^T \Sigma_{xx}^{-1} (x - \mu_x)}{2}\right)}{(2\pi)^{\frac{D_x}{2}} \sqrt{|\Sigma_{xx}|}}$$

$$F_{yy} = \frac{\exp\left(-\frac{(y - \mu_y)^T \Sigma_{yy}^{-1} (y - \mu_y)}{2}\right)}{(2\pi)^{\frac{D_y}{2}} \sqrt{|\Sigma_{yy}|}}$$


Marginals/Conditionals for Dependent Gaussians

$$p(x, y) = \frac{\exp\left(-\frac{1}{2}\left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)^T \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}^{-1} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)\right)}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}}$$




Dependent Gaussian Marginals/Conditionals

$$p(x, y) = \frac{\exp\left(-\frac{1}{2}\left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)^T \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}^{-1} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)\right)}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}}$$



Use affine transformation $\begin{bmatrix} x \\ y \end{bmatrix} \rightarrow \begin{bmatrix} x \\ w = y - \Sigma_{yx} \Sigma_{xx}^{-1} (x - \mu_x) \end{bmatrix}$

$$\text{cov}(x, w) = \Sigma_{yx} - \Sigma_{yx} \Sigma_{xx}^{-1} \Sigma_{xx} = 0$$



for normal variables, uncorrelated=independent

$$\mu_w = \mu_y$$

Gaussian Marginals/Conditionals

- Conditional & marginal from joint: $p(y|x) = \frac{p(x,y)}{p(x)} = \frac{p(x,y)}{\int_y p(x,y)}$

- Gaussian: $p(z|\mu, \Sigma) = \frac{1}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}} \exp\left(-\frac{1}{2} (z - \mu)^T \Sigma^{-1} (z - \mu)\right)$

$$p(x, y) = \frac{\exp\left(-\frac{1}{2} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)^T \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}^{-1} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)\right)}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}}$$

$$p(x) = \frac{1}{(2\pi)^{\frac{D_x}{2}} \sqrt{|\Sigma_{xx}|}} \exp\left(-\frac{1}{2} (x - \mu_x)^T \Sigma_{xx}^{-1} (x - \mu_x)\right) = N(\mu_x, \Sigma_{xx})$$

Regression with Gaussians

•Conditional & marginal from joint: $p(y|x) = \frac{p(x,y)}{p(x)} = \frac{p(x,y)}{\int_y p(x,y)}$

•Gaussian: $p(z|\mu, \Sigma) = \frac{1}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}} \exp\left(-\frac{1}{2} (z - \mu)^T \Sigma^{-1} (z - \mu)\right)$


$$p(x, y) = \frac{\exp\left(-\frac{1}{2} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)^T \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}^{-1} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}\right)\right)}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma|}}$$

$$p(x) = \frac{1}{(2\pi)^{\frac{D_x}{2}} \sqrt{|\Sigma_{xx}|}} \exp\left(-\frac{1}{2} (x - \mu_x)^T \Sigma_{xx}^{-1} (x - \mu_x)\right) = N(\mu_x, \Sigma_{xx})$$

$$p(y|x) = N(\mu_y + \Sigma_{yx} \Sigma_{xx}^{-1} (x - \mu_x), \Sigma_{yy} - \Sigma_{yx} \Sigma_{xx}^{-1} \Sigma_{xy})$$







•Here argmax is conditional expectation:

$$\hat{y} = \mu_y + \Sigma_{yx} \Sigma_{xx}^{-1} (x - \mu_x)$$




Principal Components Analysis

- Gaussians: for Classification, Regression... & Compression!
- Data can be constant in some directions, changes in others
- Use Gaussian to find directions of high/low variance
- Intuition: Regression = know x , measure noisy y
PCA = measure noisy x, y




Principal Components Analysis

- Idea: instead of writing data in all its dimensions, only write it as mean + steps along one direction




$$\begin{bmatrix} x_i \\ y_i \end{bmatrix} \approx \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix} + c_i \begin{bmatrix} v_x \\ v_y \end{bmatrix}$$



- More generally, keep a subset of dimensions

$$\vec{z}_i = \vec{\mu} + \sum_{j=1}^C c_{ij} \vec{v}_j$$



- Compression method: instead of \vec{z}_i , only save \vec{c}_i
- Optimal directions: along eigenvectors of Σ
- Which directions to keep: highest eigenvalues (new variances)

Principal Components Analysis

- If we have eigenvectors, mean and coefficients:

$$\vec{z}_i = \vec{\mu} + \sum_{j=1}^C c_{ij} \vec{v}_j$$

- Get eigenvectors $\Sigma = V \Lambda V^T$

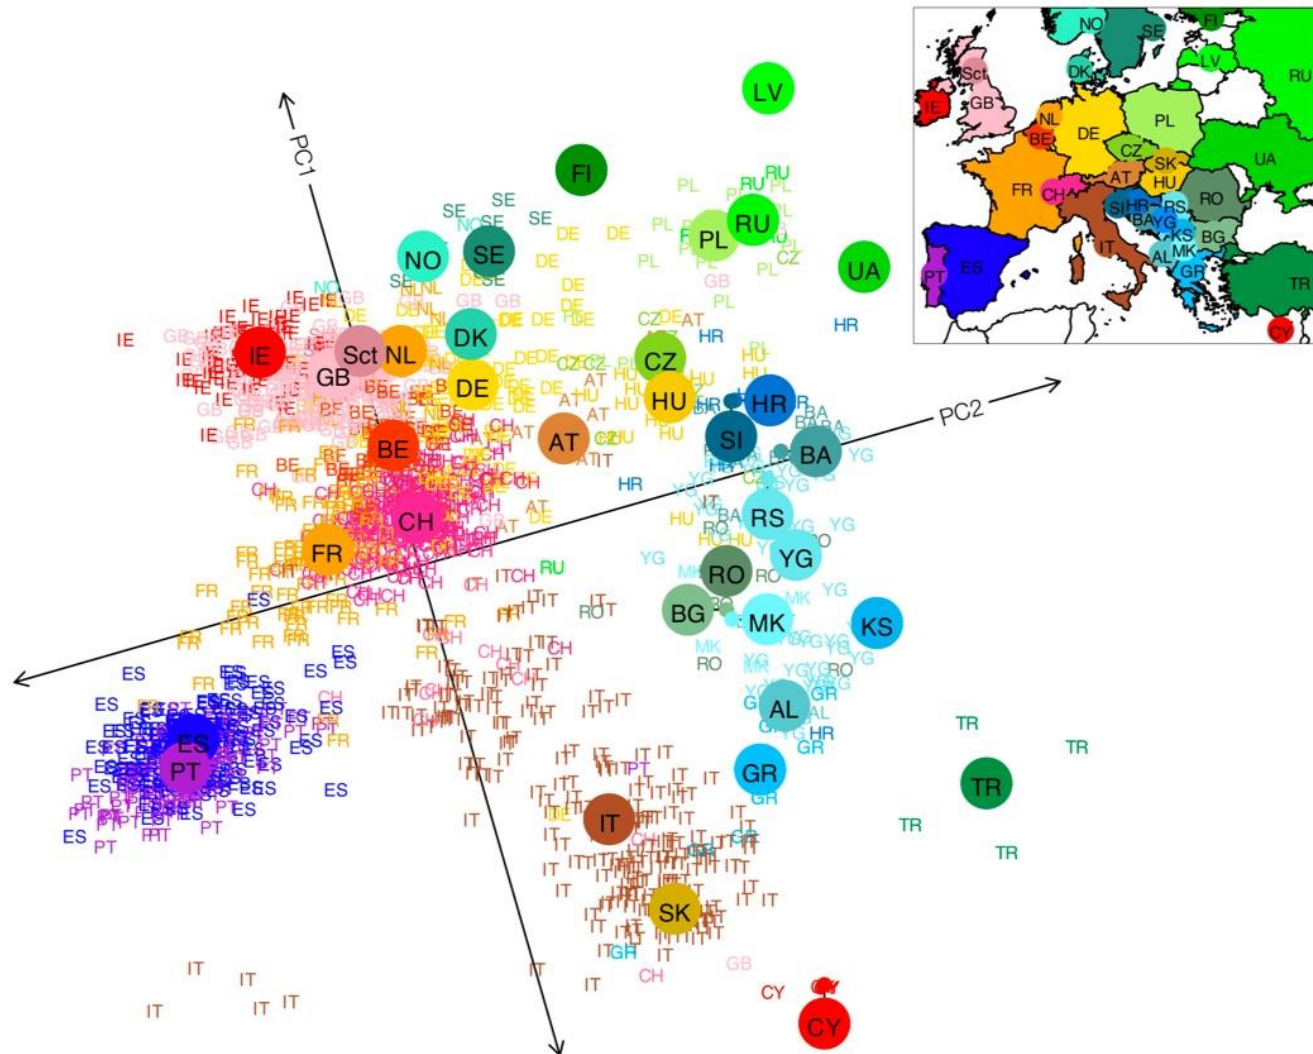
$$\begin{bmatrix} \Sigma(1,1) & \Sigma(1,2) & \Sigma(1,3) \\ \Sigma(2,1) & \Sigma(2,2) & \Sigma(2,3) \\ \Sigma(3,1) & \Sigma(3,2) & \Sigma(3,3) \end{bmatrix} = \begin{bmatrix} [\vec{v}_1] & [\vec{v}_2] & [\vec{v}_3] \end{bmatrix} \begin{bmatrix} \lambda_1 & 0 & 0 \\ 0 & \lambda_2 & 0 \\ 0 & 0 & \lambda_3 \end{bmatrix} \begin{bmatrix} [\vec{v}_1] & [\vec{v}_2] & [\vec{v}_3] \end{bmatrix}^T$$

- Eigenvectors are orthonormal: $\vec{v}_i^T \vec{v}_j = \delta_{ij}$
- In coordinates of v , Gaussian is diagonal, $\text{cov} = \Lambda$
- All eigenvalues are non-negative $\lambda_i \geq 0$
- Higher eigenvalues are higher variance, use the top C ones

$$\lambda_1 \geq \lambda_2 \geq \lambda_3 \geq \lambda_4 \geq \dots$$

- To compute the coefficients: $c_{ij} = (\vec{z}_i - \vec{\mu})^T \vec{v}_j$

PCA in Genetics




How Many Eigenvectors Needed?

- ◆ Not necessary to fully decompose
- ◆ Greedy algorithm:
 - Iteratively find largest λ & peel off vector


How Many Eigenvectors Needed?

- ◆ Not necessary to fully decompose
- ◆ Greedy algorithm:
 - Iteratively find largest λ & peel off vector
 - Repeatedly multiply a start vector & normalize
- ◆ Tracy-Widom statistics:
 - Null distribution of %variance λ_k explains

Two Gaussians





Two Unknown Gaussians



Mixtures for More Flexibility

- With mixtures (e.g. mixtures of Gaussians) we can handle more complicated (e.g. multi-bump, nonlinear) distributions.

subpopulations: G1=compact car
 G2=mid-size car
 G3=cadillac




- In fact, if we have enough Gaussians (maybe infinite) we can approximate any distribution...


Mixtures as Hidden Variables

- Consider a dataset with K subpopulations but don't know which subpopulation each point belongs to

e.g. looking at height of adult people, we see $K=2$ subpopulations: males & females



e.g. looking at pitch and height of people we see $K=2$ subpopulations: males & females




- Because of the 'hidden' variable (y can be 1 or 2), these distributions are not Gaussians but **Mixture of Gaussians**

$$p(\vec{x}) =$$


Mixtures as Hidden Variables

- Consider a dataset with K subpopulations but don't know which subpopulation each point belongs to

e.g. looking at height of adult people, we see $K=2$ subpopulations: males & females



e.g. looking at pitch and height of people we see $K=2$ subpopulations: males & females




- Because of the 'hidden' variable (y can be 1 or 2), these distributions are not Gaussians but **Mixture of Gaussians**

$$\begin{aligned}
 p(\vec{x}) &= \sum_y p(y)p(\vec{x}|y) = \sum_y \pi_y N(\vec{x}|\vec{\mu}_y, \Sigma_y) \\
 &= \sum_y \pi_y \frac{1}{(2\pi)^{\frac{D}{2}} \sqrt{|\Sigma_y|}} \exp\left(-\frac{1}{2}(\vec{x} - \vec{\mu}_y)^T \Sigma_y^{-1}(\vec{x} - \vec{\mu}_y)\right)
 \end{aligned}$$

Hidden / Unlabeled = Clustering

- Recall classification problem:
maximize the log-likelihood of
data given models:

$l =$



Hidden / Unlabeled = Clustering

- Recall classification problem:



maximize the log-likelihood of data given models:

$$l = \sum_{n=1}^N \log p(\vec{x}_n, y_n | \pi, \mu, \Sigma) \\ = \sum_{n=1}^N \log \pi_{y_n} + \log N(\vec{x}_n | \vec{\mu}_{y_n}, \Sigma_{y_n})$$

- If we don't know the class
treat it as a hidden variable

maximize the log-likelihood with unlabeled data:

$$l = \sum_{n=1}^N \log p(\vec{x}_n | \pi, \mu, \Sigma) =$$



Hidden / Unlabeled = Clustering

- Recall classification problem:

maximize the log-likelihood of data given models:



$$l = \sum_{n=1}^N \log p(\vec{x}_n, y_n | \pi, \mu, \Sigma) \\ = \sum_{n=1}^N \log \pi_{y_n} + \log N(\vec{x}_n | \vec{\mu}_{y_n}, \Sigma_{y_n})$$

- If we don't know the class treat it as a hidden variable

maximize the log-likelihood with unlabeled data:

$$l = \sum_{n=1}^N \log p(\vec{x}_n | \pi, \mu, \Sigma) = \sum_{n=1}^N \log \sum_{y=1}^K p(\vec{x}_n, y | \pi, \mu, \Sigma) \\ = \sum_{n=1}^N \log(\pi_1 N(\vec{x}_n | \vec{\mu}_1, \Sigma_1) + \dots + \pi_K N(\vec{x}_n | \vec{\mu}_K, \Sigma_K))$$

- Instead of classification, we now have a **clustering** problem



Hidden / Unlabeled = Clustering


- Represent each hidden y integer (1 to K) with a hidden binary indicator vector z

$$\vec{z} \in \mathbf{B}^K, \sum_{i=1}^K \vec{z}(i) = 1 \text{ or}$$

$$\vec{z} \in \{\vec{\delta}_1, \dots, \vec{\delta}_K\} \text{ where } \vec{\delta}_i(i) = 1, \vec{\delta}_i(j) = 0 \text{ for } i \neq j$$

- Each likelihood requires summing over the possible z

$$p(\vec{x}|\theta) =$$



Hidden / Unlabeled = Clustering


- Represent each hidden y integer (1 to K) with a hidden binary indicator vector z

$$\vec{z} \in \mathbf{B}^K, \sum_{i=1}^K \vec{z}(i) = 1 \text{ or}$$

$$\vec{z} \in \{\vec{\delta}_1, \dots, \vec{\delta}_K\} \text{ where } \vec{\delta}_i(i) = 1, \vec{\delta}_i(j) = 0 \text{ for } i \neq j$$

- Each likelihood requires summing over the possible z

$$p(\vec{x}|\theta) = \sum_z p(\vec{z}|\theta)p(\vec{x}|\vec{z}, \theta) = \sum_{i=1}^K p(\vec{z} = \vec{\delta}_i|\theta)p(\vec{x}|\vec{z} = \vec{\delta}_i, \theta)$$



Hidden / Unlabeled = Clustering

- Represent each hidden y integer (1 to K) with a hidden binary indicator vector z

$$\vec{z} \in \mathbf{B}^K, \sum_{i=1}^K \vec{z}(i) = 1 \text{ or}$$

$$\vec{z} \in \{\vec{\delta}_1, \dots, \vec{\delta}_K\} \text{ where } \vec{\delta}_i(i) = 1, \vec{\delta}_i(j) = 0 \text{ for } i \neq j$$

- Each likelihood requires summing over the possible z


$$p(\vec{x}|\theta) = \sum_z p(\vec{z}|\theta)p(\vec{x}|\vec{z}, \theta) = \sum_{i=1}^K p(\vec{z} = \vec{\delta}_i|\theta)p(\vec{x}|\vec{z} = \vec{\delta}_i, \theta)$$

mixing proportions (prior) $= \pi = p(\vec{z} = \vec{\delta}_i|\pi)$

mixture components $= p(\vec{x}|\vec{z} = \vec{\delta}_i, \theta)$

posteriors (responsibilities) $= \tau_{n,i} =$

log likelihood $=$



Hidden / Unlabeled = Clustering

- Represent each hidden y integer (1 to K) with a hidden binary indicator vector \vec{z}

$$\vec{z} \in \mathbf{B}^K, \sum_{i=1}^K \vec{z}(i) = 1 \text{ or}$$

$$\vec{z} \in \{\vec{\delta}_1, \dots, \vec{\delta}_K\} \text{ where } \vec{\delta}_i(i) = 1, \vec{\delta}_i(j) = 0 \text{ for } i \neq j$$

- Each likelihood requires summing over the possible \vec{z}

$$p(\vec{x}|\theta) = \sum_{\vec{z}} p(\vec{z}|\theta) p(\vec{x}|\vec{z}, \theta) = \sum_{i=1}^K p(\vec{z} = \vec{\delta}_i|\theta) p(\vec{x}|\vec{z} = \vec{\delta}_i, \theta)$$


mixing proportions (prior) $= \pi = p(\vec{z} = \vec{\delta}_i|\pi)$

mixture components $= p(\vec{x}|\vec{z} = \vec{\delta}_i, \theta)$


posteriors (responsibilities) $= \tau_{n,i} = p(\vec{z} = \vec{\delta}_i|\vec{x}_n, \theta) = \frac{p(\vec{x}_n|\vec{z}=\vec{\delta}_i, \theta) p(\vec{z}=\vec{\delta}_i|\theta)}{p(\vec{x}_n|\theta)}$

log likelihood $= \sum_{n=1}^N \log p(\vec{x}_n|\pi, \mu, \Sigma) = \sum_{n=1}^N \log \sum_{i=1}^K \pi_i p(\vec{x}_n|\vec{\mu}_i, \Sigma_i)$

- Can't easily take derivatives of log-likelihood and set to 0.
- Not nice, seems to need gradient ascent...
- Or, can we break up mixture into smaller Gaussian steps?

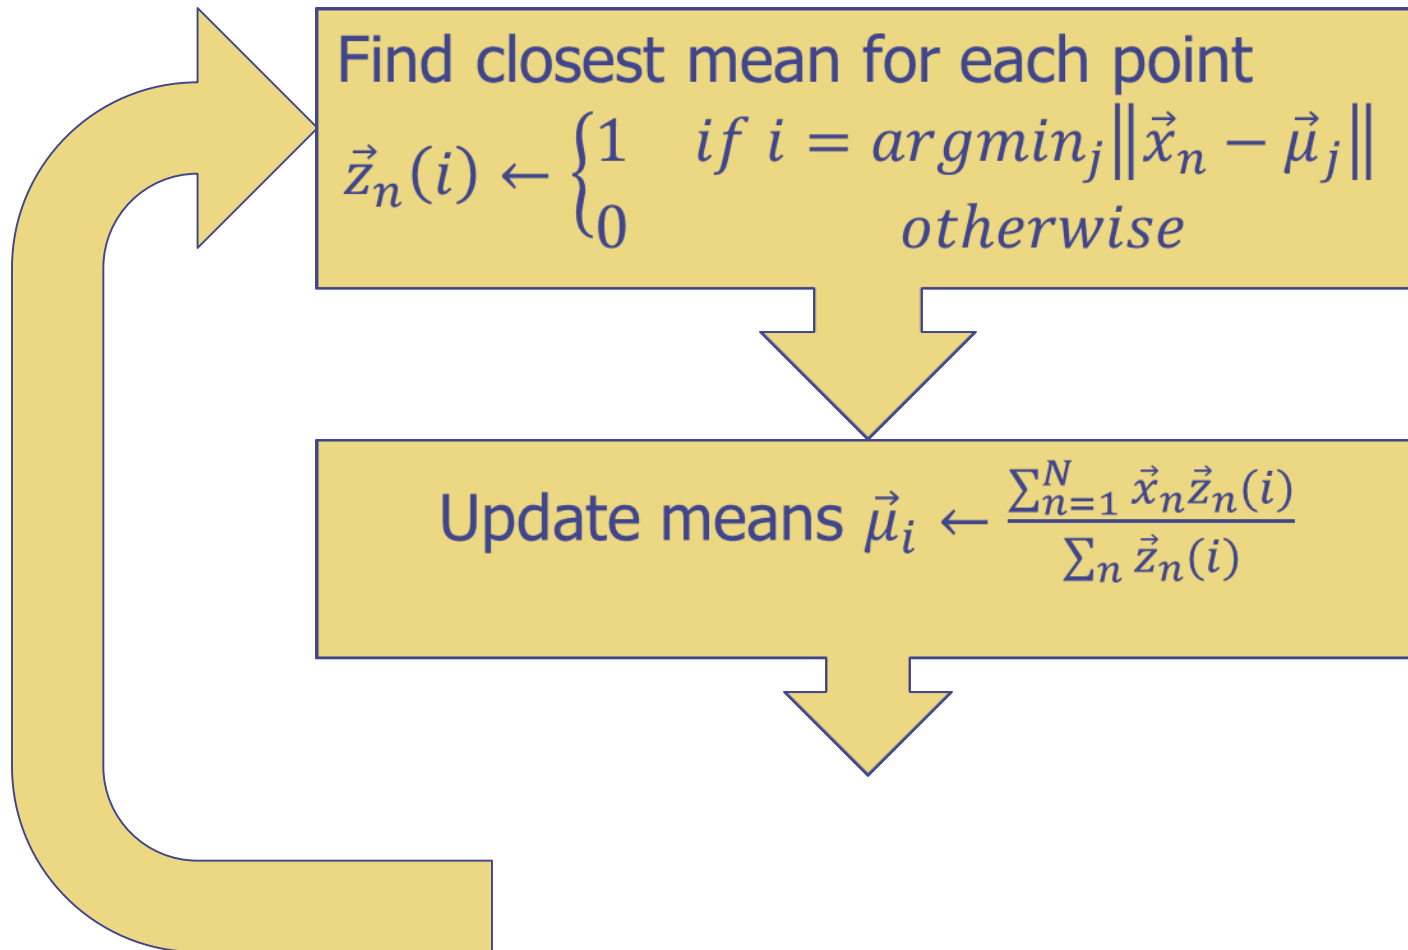


K-Means Clustering

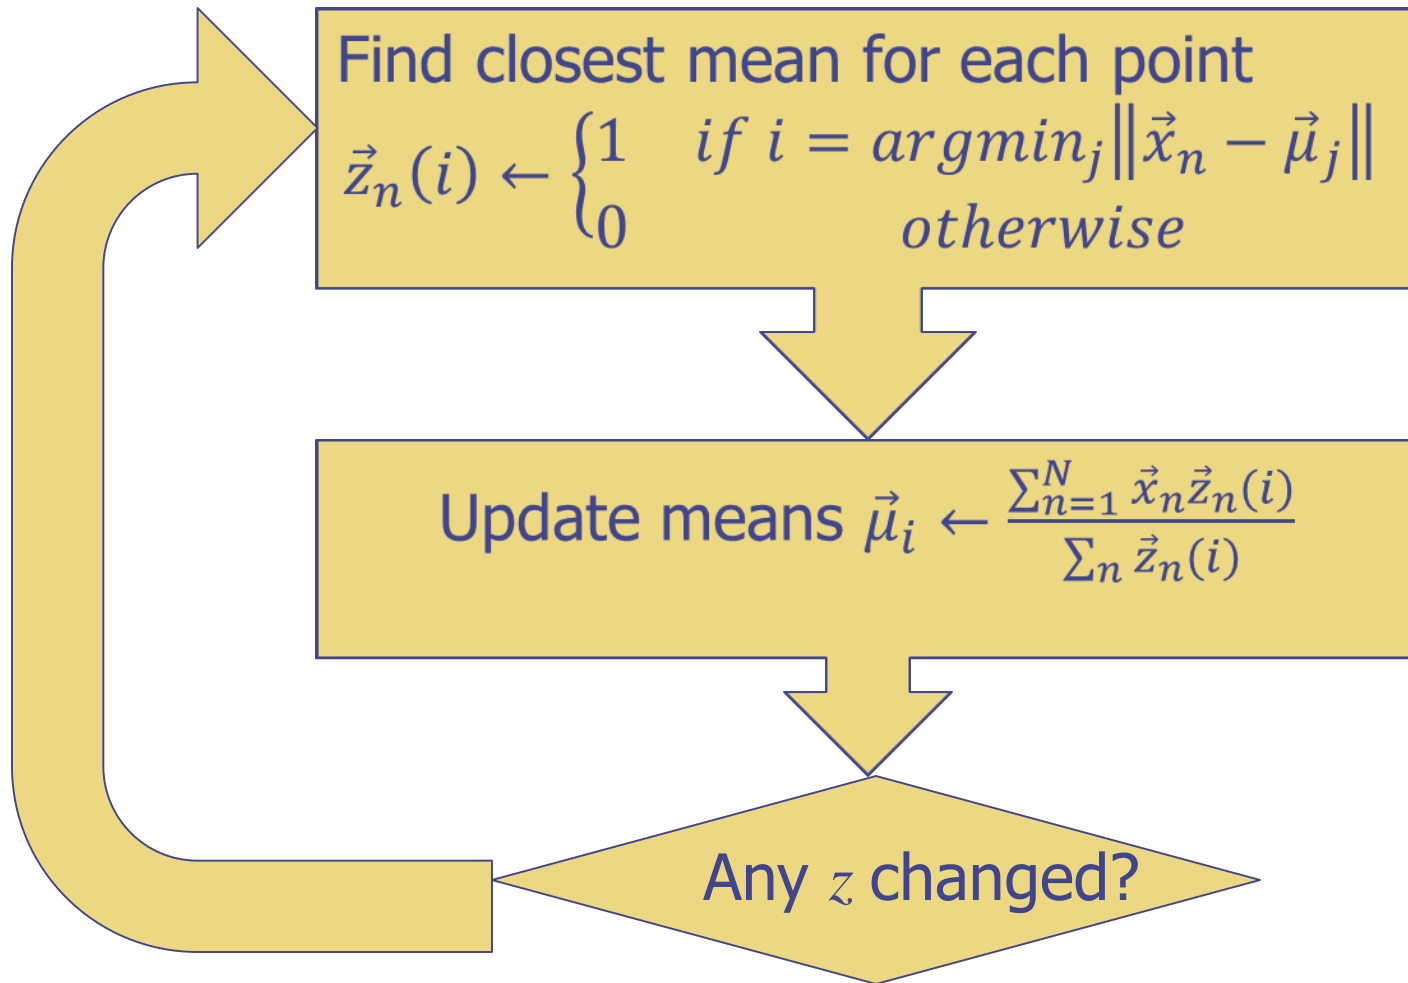


- An old “heuristic” clustering algorithm
- Gobble up data with a divide & conquer scheme
- Assume each point x has a discrete multinomial vector z
- Chicken and Egg problem:
 - If know classes, we can get model (max likelihood!)
 - If know the model, we can predict the classes (classifier!)


The K-Means Algorithm




The K-Means Algorithm



The K-Means Algorithm



K-Means Clustering




- Geometric, each point goes to closest Gaussian
- Recompute the means by their assigned points
- Minimizing $\min_{\mu} \min_{z} J(\vec{\mu}_1, \dots, \vec{\mu}_K, \vec{z}_1, \dots, \vec{z}_N)$ cost function:

$$J(\vec{\mu}_1, \dots, \vec{\mu}_K, \vec{z}_1, \dots, \vec{z}_N) = \sum_{n=1}^N \sum_{i=1}^K \vec{z}_n(i) \|\vec{x}_n - \vec{\mu}_i\|^2$$

$$\vec{z}_n(i) = \begin{cases} 1 & \text{if } i = \operatorname{argmin}_j \|\vec{x}_n - \vec{\mu}_j\| \\ 0 & \text{otherwise} \end{cases} \quad \text{and} \quad \vec{\mu}_i = \frac{\sum_{n=1}^N \vec{x}_n \vec{z}_n(i)}{\sum_n \vec{z}_n(i)}$$

K-Means Clustering




- Geometric, each point goes to closest Gaussian
- Recompute the means by their assigned points
- Minimizing $\min_{\mu} \min_z J(\vec{\mu}_1, \dots, \vec{\mu}_K, \vec{z}_1, \dots, \vec{z}_N)$ cost function:

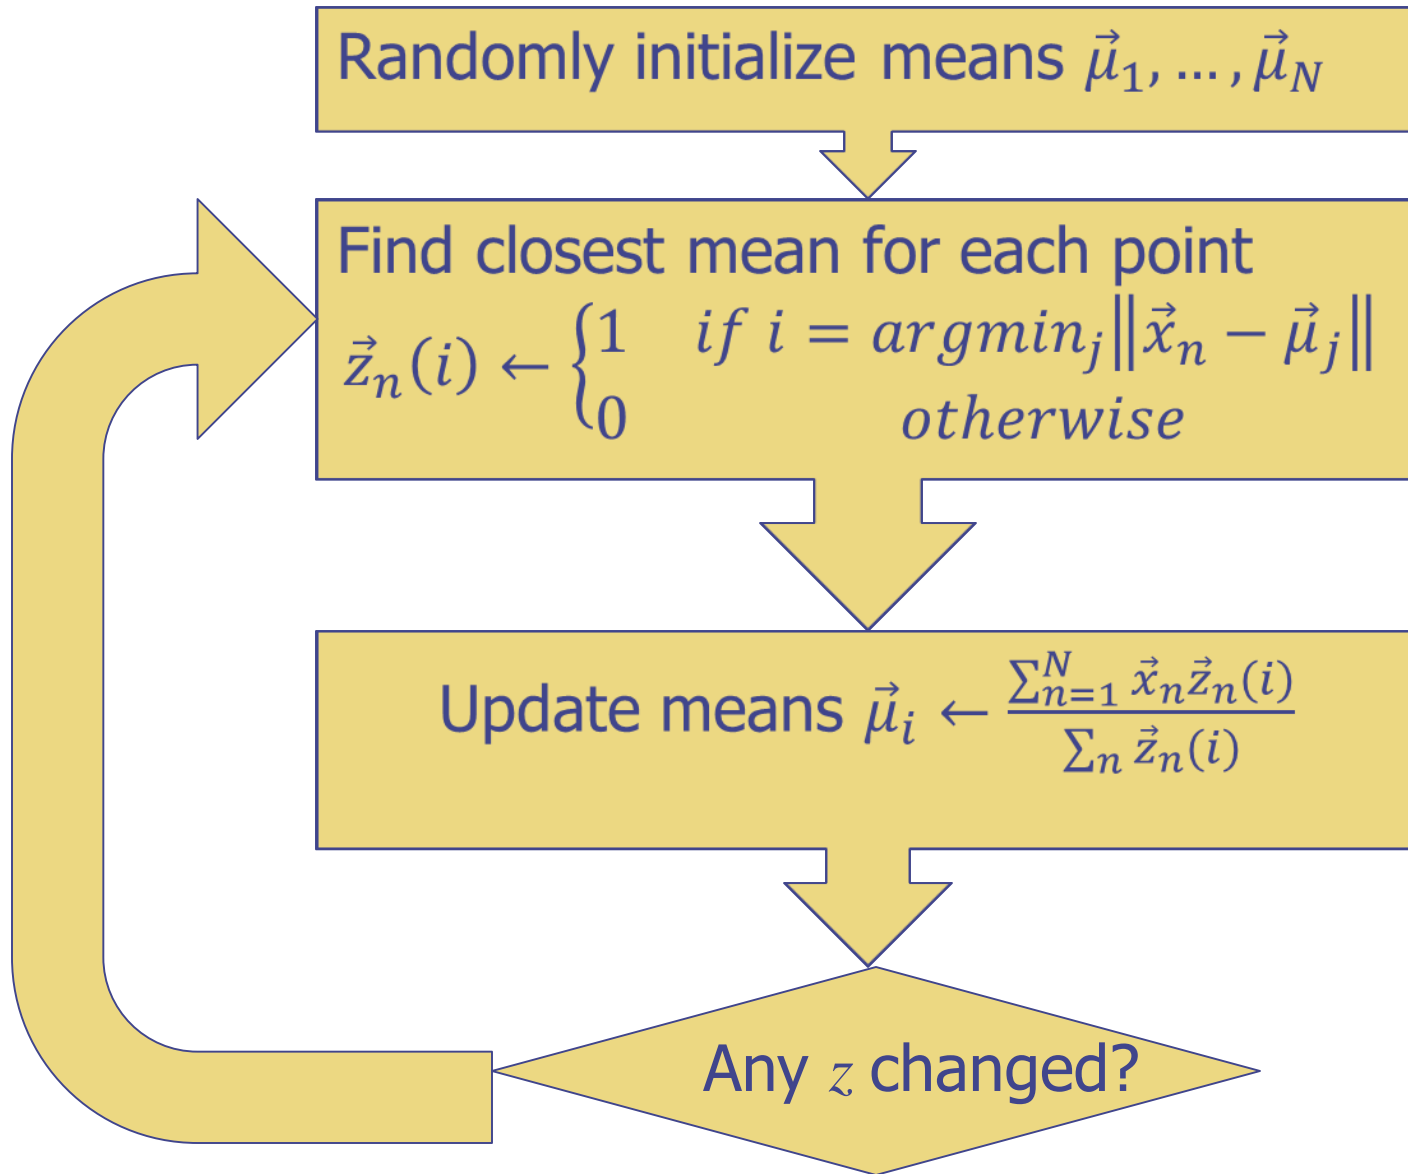
$$J(\vec{\mu}_1, \dots, \vec{\mu}_K, \vec{z}_1, \dots, \vec{z}_N) = \sum_{n=1}^N \sum_{i=1}^K \vec{z}_n(i) \|\vec{x}_n - \vec{\mu}_i\|^2$$

$$\vec{z}_n(i) = \begin{cases} 1 & \text{if } i = \operatorname{argmin}_j \|\vec{x}_n - \vec{\mu}_j\| \\ 0 & \text{otherwise} \end{cases} \quad \text{and} \quad \vec{\mu}_i = \frac{\sum_{n=1}^N \vec{x}_n \vec{z}_n(i)}{\sum_n \vec{z}_n(i)}$$

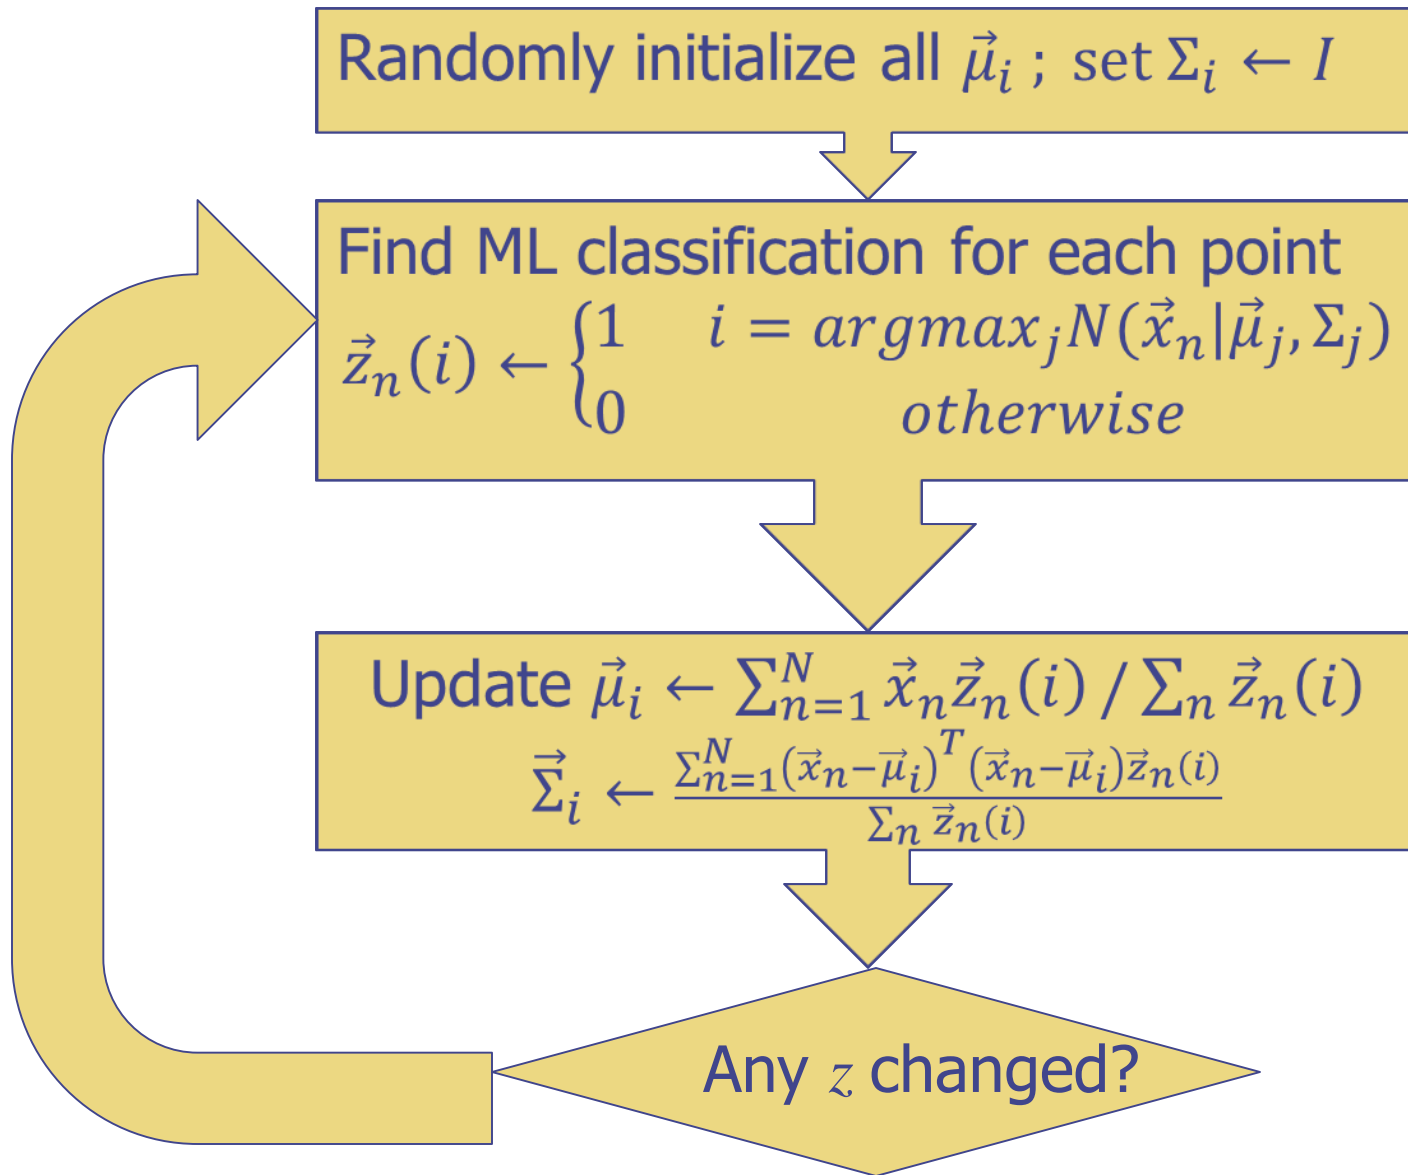
- Guaranteed to improve per iteration and converge
- Like **Coordinate Descent** (lock one var, maximize the other)
- A.k.a. **Axis-Parallel Optimization** or **Alternating Minimization**




The K-Means Algorithm



The Gaussian K-Means Algorithm



K-Means Clustering



- Geometric, each point goes to closest Gaussian
- Recompute the means by their assigned points
- Minimizing $\min_{\mu, \Sigma} \min_z J(\vec{\mu}_1, \dots, \vec{\mu}_K, \Sigma_1, \dots, \Sigma_K, \vec{z}_1, \dots, \vec{z}_N)$ cost:

$$J(\vec{\mu}_1, \dots, \vec{\mu}_K, \vec{z}_1, \dots, \vec{z}_N) = \sum_{n=1}^N \sum_{i=1}^K \vec{z}_n(i) (\vec{x}_n - \vec{\mu}_i)^T \Sigma^{-1} (\vec{x}_n - \vec{\mu}_i)$$

$$\vec{z}_n(i) = \begin{cases} 1 & \text{if } i = \operatorname{argmin}_j (\vec{x}_n - \vec{\mu}_j)^T \Sigma^{-1} (\vec{x}_n - \vec{\mu}_j) \\ 0 & \text{otherwise} \end{cases}$$

- Guaranteed to improve per iteration and converge
- Like **Coordinate Descent** (lock one var, maximize the other)
- A.k.a. **Axis-Parallel Optimization** or **Alternating Minimization**

